

Mean Returns in Time Series - Restarting after NA values - rstudio

Question

Has anyone encountered calculating historical mean log returns in time series datasets?

The dataset is ordered by individual security first and by time for each respective security. I am trying to form a historical mean log return, i.e. the mean log return for the security from its first appearance in the dataset to date, for each point in time for each security.

Luckily, the return time series contains NAs between returns for differing securities. My idea is to calculate a historical mean that restarts after each NA that appears.

A simple `cumsum()` probably will not do it, as the NAs will have to be dropped.

I thought about using `rollmean()`, if I only knew an efficient way to specify the 'width' parameter to the length of the vector of consecutive preceding non-NAs. The current approach I am taking, based on Count how many consecutive values are true, takes significantly too much time, given the size of the data set I am working with. For any `x` of the form `x : [r(1) r(2) ... r(N)]`, where `r(2)` is the log return in period 2:

```
df <- data.frame(x, zcount = NA)
df[1,2] = 0 #df$x[1]=NA by construction of the data set
for(i in 2:nrow(df))
df$zcount[i] <- ifelse(!is.na(df$x[i]), df$zcount[i-1]+1, 0)
```

Any idea how to speed this up would be highly appreciated!

Answer

You will need to reshape the data.frame to apply the `cumsum` function over each security. Here's how:

First, I'll generate some data on 100 securities over 100 months which I think corresponds to your description of the data set

```
securities <- 100
months <- 100

time <- seq.Date(as.Date("2010/1/1"), by = "months", length.out = months)
ID <- rep(paste0("sec", 1:months), each = securities)
returns <- rnorm(securities * months, mean = 0.08, sd = 2)

df <- data.frame(time, ID, returns)

head(df)
```

```
##      time    ID    returns
## 1 2010-01-01 sec1 -1.49709654
## 2 2010-02-01 sec1  2.51423770
## 3 2010-03-01 sec1 -0.02174069
## 4 2010-04-01 sec1  0.15159940
## 5 2010-05-01 sec1 -4.69473924
## 6 2010-06-01 sec1 -3.15394053
```

```
tail(df)
```

```
##           time      ID    returns
## 9995 2017-11-01 sec100 -1.6294895
## 9996 2017-12-01 sec100  1.0715943
## 9997 2018-01-01 sec100 -0.1815839
## 9998 2018-02-01 sec100  0.4914217
## 9999 2018-03-01 sec100 -1.9122788
## 10000 2018-04-01 sec100  2.3707944
```

Now, you must reshape your data so that each security column contains its returns, and each row represents the date.

```
library(tidyr)
df_wide <- spread(df, ID, returns)

head(df_wide[,1:6])
```

```
##           time      sec1      sec10      sec100      sec11      sec12
## 1 2010-01-01 -1.49709654 -0.4024286 -3.998970 -3.4819828  3.7054058
## 2 2010-02-01  2.51423770 -2.6837960 -1.446198 -2.6753800  2.9874788
## 3 2010-03-01 -0.02174069 -0.8093903 -2.729755 -1.7519346 -1.4416595
## 4 2010-04-01  0.15159940 -0.3487980 -2.713428 -3.1053596 -1.0103877
## 5 2010-05-01 -4.69473924 -0.6979461  3.276191  1.2752365 -0.4962654
## 6 2010-06-01 -3.15394053 -0.3028585 -3.260437  0.1449569 -4.2148564
```

Once this is done, you can use the `apply` function to `sum` every column which now represents each security. Or use the `cumsum` function. Notice the data object `df_wide[-1]`, which drops the time column. This is necessary to avoid the sum or cumsum functions throwing an error.

```
matrix_sum <- apply(df_wide[-1], 2, FUN = sum)

matrix_cumsum <- apply(df_wide[-1], 2, FUN = cumsum)
```

Now, add the time column back as a `data.frame` if you like:

```
df_final <- data.frame(time = df_wide[,1], matrix_cumsum)

head(df_final[,1:6])
```

```
##           time      sec1      sec10      sec100      sec11      sec12
## 1 2010-01-01 -1.4970965 -0.4024286 -3.998970 -3.481983  3.7054058
## 2 2010-02-01  1.0171412 -3.0862246 -5.445169 -6.157363  6.6928846
## 3 2010-03-01  0.9954005 -3.8956149 -8.174924 -7.909297  5.2512251
## 4 2010-04-01  1.1469999 -4.2444129 -10.888352 -11.014657  4.2408374
## 5 2010-05-01 -3.5477394 -4.9423590 -7.612161 -9.739421  3.7445720
## 6 2010-06-01 -6.7016799 -5.2452175 -10.872599 -9.594464 -0.4702844
```

```
tail(df_final[,1:6])
```

```
##           time      sec1      sec10      sec100      sec11      sec12
## 95 2017-11-01 14.06059  2.2175772 -22.67810 -20.40392 34.74430
## 96 2017-12-01 13.47993  0.3382542 -21.60651 -19.39266 33.59920
## 97 2018-01-01 14.94843  1.6157574 -21.78809 -18.57979 33.57235
## 98 2018-02-01 15.73950 -0.9858401 -21.29667 -21.45400 32.18920
## 99 2018-03-01 16.14455  0.2420269 -23.20895 -23.06486 28.29411
## 100 2018-04-01 20.28924  1.1684538 -20.83815 -26.20104 27.34003
```